Chapter 18. with, match, and else Blocks

Context managers may end up being almost as important as the subroutine itself. We've only scratched the surface with them. [...] Basic has a `with` statement, there are `with` statements in lots of languages. But they don't do the same thing, they all do something very shallow, they save you from repeated dotted [attribute] lookups, they don't do setup and tear down. Just because it's the same name don't think it's the same thing. The `with` statement is a very big deal.[1]

—Raymond Hettinger, eloquent Python evangelist

This chapter is about control flow features that are not so common in other languages, and for this reason tend to be overlooked or underused in Python. They are:

- The `with` statement and context manager protocol
- Pattern matching with `match/case`
- The `else` clause in `for`, `while`, and `try` statements

The `with` statement sets up a temporary context and reliably tears it down, under the control of a context manager object. This prevents errors and reduces boilerplate code, making APIs at the same time safer and easier to use. Python programmers are finding lots of uses for `with` blocks beyond automatic file closing.

We've seen pattern matching in previous chapters, but here we'll see how the grammar of a language can be expressed as sequence patterns. That observation explains why `match/case` is an effective tool to create language processors that are easy to understand and extend. We'll study a complete interpreter for a small but functional subset of the Scheme language. The same ideas could be applied to develop a template language or a DSL (Domain-Specific Language) to encode business rules in a larger system.

The `else` clause is not a big deal, but it does help convey intention when properly used together with `for`, `while`, and `try`.

What's New in This Chapter

["Pattern Matching in lis.py: A Case Study"](#) is a new section.

I updated ["The contextlib Utilities"](#) to cover a few features of the `contextlib` module added since Python 3.6, and the new parenthesized context managers syntax introduced in Python 3.10.

Let's start with the powerful `with` statement.

Context Managers and with Blocks

Context manager objects exist to control a `with` statement, just like iterators exist to control a `for` statement.

The `with` statement was designed to simplify some common uses of `try/finally`, which guarantees that some operation is performed after a block of code, even if the block is terminated by `return`, an exception, or a `sys.exit()` call. The code in the `finally` clause usually releases a critical resource or restores some previous state that was temporarily changed.

The Python community is finding new, creative uses for context managers. Some examples from the standard library are:

- Managing transactions in the `sqlite3` module—see ["Using the connection as a context manager"](#).
- Safely handling locks, conditions, and semaphores—as described in the `threading` [module documentation](#).
- Setting up custom environments for arithmetic operations with `Decimal` objects—see the [`decimal.localcontext` documentation](#).
- Patching objects for testing—see the [`unittest.mock.patch` function](#).

The context manager interface consists of the `__enter__` and `__exit__` methods. At the top of the `with`, Python calls the `__enter__` method of the context manager object. When the `with` block completes or terminates for any reason, Python calls `__exit__` on the context manager object.

The most common example is making sure a file object is closed. Example 18-1 is a detailed demonstration of using `with` to close a file.

Example 18-1. Demonstration of a file object as a context manager

```
>>> with open('mirror.py') as fp:  ❶
...     src = fp.read(60)  ❷
...
>>> len(src)
60
>>> fp  ❸
<_io.TextIOWrapper name='mirror.py' mode='r' encoding='UTF-8'>
>>> fp.closed, fp.encoding  ❹
(True, 'UTF-8')
>>> fp.read(60)  ❺
Traceback (most recent call last):
  File "<stdin>", line 1, in <module>
ValueError: I/O operation on closed file.
```

❶ `fp` is bound to the opened text file because the file's `__enter__` method returns `self`.

❷ Read `60` Unicode characters from `fp`.

❸ The `fp` variable is still available— `with` blocks don't define a new scope, as functions do.

❹ We can read the attributes of the `fp` object.

❺ But we can't read more text from `fp` because at the end of the `with` block, the `TextIOWrapper.__exit__` method was called, and it closed the file.

The first callout in Example 18-1 makes a subtle but crucial point: the context manager object is the result of evaluating the expression after

`with` , but the value bound to the target variable (in the `as` clause) is the result returned by the `__enter__` method of the context manager object.

It just happens that the `open()` function returns an instance of `TextIOWrapper` , and its `__enter__` method returns `self` . But in a different class, the `__enter__` method may also return some other object instead of the context manager instance.

When control flow exits the `with` block in any way, the `__exit__` method is invoked on the context manager object, not on whatever was returned by `__enter__` .

The `as` clause of the `with` statement is optional. In the case of `open` , we always need it to get a reference to the file, so that we can call methods on it. But some context managers return `None` because they have no useful object to give back to the user.

Example 18-2 shows the operation of a perfectly frivolous context manager designed to highlight the distinction between the context manager and the object returned by its `__enter__` method.

Example 18-2. Test-driving the `LookingGlass` context manager class

```
>>> from mirror import LookingGlass
>>> with LookingGlass() as what:    ❶
...        print('Alice, Kitty and Snowdrop')   ❷
...        print(what)
...
pordwonS dna yttiK ,ecilA
YKCOWREBBAJ
>>> what    ❸
'JABBERWOCKY'
>>> print('Back to normal.')    ❹
Back to normal.
```

❶ The context manager is an instance of `LookingGlass` ; Python calls `__enter__` on the context manager and the result is bound to `what` .

❷ Print a `str`, then the value of the target variable `what`. The output of each `print` will come out reversed.

❸ Now the `with` block is over. We can see that the value returned by `__enter__`, held in `what`, is the string `'JABBERWOCKY'`.

❹ Program output is no longer reversed.

Example 18-3 shows the implementation of `LookingGlass`.

Example 18-3. mirror.py: code for the `LookingGlass` context manager class

```python
import sys

class LookingGlass:

    def __enter__(self):  ❶
        self.original_write = sys.stdout.write  ❷
        sys.stdout.write = self.reverse_write  ❸
        return 'JABBERWOCKY'  ❹

    def reverse_write(self, text):  ❺
        self.original_write(text[::-1])

    def __exit__(self, exc_type, exc_value, traceback):  ❻
        sys.stdout.write = self.original_write  ❼
        if exc_type is ZeroDivisionError:  ❽
            print('Please DO NOT divide by zero!')
            return True  ❾
        ❿
```

❶ Python invokes `__enter__` with no arguments besides `self`.

❷ Hold the original `sys.stdout.write` method, so we can restore it later.

❸ Monkey-patch `sys.stdout.write`, replacing it with our own method.

❹

Return the `'JABBERWOCKY'` string just so we have something to put in the target variable `what`.

❺ Our replacement to `sys.stdout.write` reverses the `text` argument and calls the original implementation.

❻ Python calls `__exit__` with `None, None, None` if all went well; if an exception is raised, the three arguments get the exception data, as described after this example.

❼ Restore the original method to `sys.stdout.write`.

❽ If the exception is not `None` and its type is `ZeroDivisionError`, print a message...

❾ ...and return `True` to tell the interpreter that the exception was handled.

❿ If `__exit__` returns `None` or any *falsy* value, any exception raised in the `with` block will be propagated.

TIP

When real applications take over standard output, they often want to replace `sys.stdout` with another file-like object for a while, then switch back to the original. The `contextlib.redirect_stdout` context manager does exactly that: just pass it the file-like object that will stand in for `sys.stdout`.

The interpreter calls the `__enter__` method with no arguments—beyond the implicit `self`. The three arguments passed to `__exit__` are:

exc_type

The exception class (e.g., `ZeroDivisionError`).

exc_value

The exception instance. Sometimes, parameters passed to the exception constructor—such as the error message—can be found in `exc_value.args`.

traceback

A `traceback` object.[2]

For a detailed look at how a context manager works, see [Example 18-4](#), where `LookingGlass` is used outside of a `with` block, so we can manually call its `__enter__` and `__exit__` methods.

Example 18-4. Exercising `LookingGlass` **without a** `with` **block**

```
>>> from mirror import LookingGlass
>>> manager = LookingGlass()   ❶
>>> manager   # doctest: +ELLIPSIS
<mirror.LookingGlass object at 0x...>
>>> monster = manager.__enter__()   ❷
>>> monster == 'JABBERWOCKY'   ❸
eurT
>>> monster
'YKCOWREBBAJ'
>>> manager   # doctest: +ELLIPSIS
>... ta tcejbo ssalGgnikooL.rorrim<
>>> manager.__exit__(None, None, None)   ❹
>>> monster
'JABBERWOCKY'
```

❶ Instantiate and inspect the `manager` instance.

❷ Call the manager's `__enter__` method and store result in `monster`.

❸ `monster` is the string `'JABBERWOCKY'`. The `True` identifier appears reversed because all output via `stdout` goes through the `write` method we patched in `__enter__`.

❹ Call `manager.__exit__` to restore the previous `stdout.write`.

Python 3.10 adopted [a new, more powerful parser](#), allowing new syntax beyond what was possible with the older [LL(1) parser](#). One syntax enhancement was to allow parenthesized context managers, like this:

```
with (
    CtxManager1() as example1,
    CtxManager2() as example2,
    CtxManager3() as example3,
):
    ...
```

Prior to 3.10, we'd have to write that as nested `with` blocks.

The standard library includes the `contextlib` package with handy functions, classes, and decorators for building, combining, and using context managers.

The contextlib Utilities

Before rolling your own context manager classes, take a look at [`contextlib`—"Utilities for `with`-statement contexts"](#) in the Python documentation. Maybe what you are about to build already exists, or there is a class or some callable that will make your job easier.

Besides the `redirect_stdout` context manager mentioned right after [Example 18-3](#), `redirect_stderr` was added in Python 3.5—it does the same as the former, but for output directed to `stderr`.

The `contextlib` package also includes:

closing

A function to build context managers out of objects that provide a `close()` method but don't implement the `__enter__`/`__exit__` interface.

suppress

A context manager to temporarily ignore exceptions given as arguments.

nullcontext

A context manager that does nothing, to simplify conditional logic around objects that may not implement a suitable context manager. It serves as a stand-in when conditional code before the `with` block may or may not provide a context manager for the `with` statement—added in Python 3.7.

The `contextlib` module provides classes and a decorator that are more widely applicable than the decorators just mentioned:

@contextmanager

A decorator that lets you build a context manager from a simple generator function, instead of creating a class and implementing the interface. See ["Using @contextmanager"](#).

AbstractContextManager

An ABC that formalizes the context manager interface, and makes it a bit easier to create context manager classes by subclassing—added in Python 3.6.

ContextDecorator

A base class for defining class-based context managers that can also be used as function decorators, running the entire function within a managed context.

ExitStack

A context manager that lets you enter a variable number of context managers. When the `with` block ends, `ExitStack` calls the stacked context managers' `__exit__` methods in LIFO order (last entered, first exited). Use this class when you don't know beforehand how many context managers you need to enter in your `with` block; for example, when opening all files from an arbitrary list of files at the same time.

With Python 3.7, `contextlib` added `AbstractAsyncContextManager`, `@asynccontextmanager`, and `AsyncExitStack`. They are similar to the equivalent utilities without the `async` part of the name, but designed for use with the new `async with` statement, covered in [Chapter 21](#).

The most widely used of these utilities is the `@contextmanager` decorator, so it deserves more attention. That decorator is also interesting because it shows a use for the `yield` statement unrelated to iteration.

Using @contextmanager

The `@contextmanager` decorator is an elegant and practical tool that brings together three distinctive Python features: a function decorator, a generator, and the `with` statement.

Using `@contextmanager` reduces the boilerplate of creating a context manager: instead of writing a whole class with `__enter__`/`__exit__` methods, you just implement a generator with a single `yield` that should produce whatever you want the `__enter__` method to return.

In a generator decorated with `@contextmanager`, `yield` splits the body of the function in two parts: everything before the `yield` will be executed at the beginning of the `with` block when the interpreter calls `__enter__`; the code after `yield` will run when `__exit__` is called at the end of the block.

[Example 18-5](#) replaces the `LookingGlass` class from [Example 18-3](#) with a generator function.

Example 18-5. mirror_gen.py: a context manager implemented with a generator

```
import contextlib
import sys

@contextlib.contextmanager   ❶
def looking_glass():
    original_write = sys.stdout.write   ❷

    def reverse_write(text):   ❸
```

```
            original_write(text[::-1])

        sys.stdout.write = reverse_write      ❹
        yield 'JABBERWOCKY'      ❺
        sys.stdout.write = original_write      ❻
```

❶ Apply the `contextmanager` decorator.

❷ Preserve the original `sys.stdout.write` method.

❸ `reverse_write` can call `original_write` later because it is available in its closure.

❹ Replace `sys.stdout.write` with `reverse_write`.

❺ Yield the value that will be bound to the target variable in the `as` clause of the `with` statement. The generator pauses at this point while the body of the `with` executes.

❻ When control exits the `with` block, execution continues after the `yield`; here the original `sys.stdout.write` is restored.

Example 18-6 shows the `looking_glass` function in operation.

Example 18-6. Test-driving the `looking_glass` context manager function

```
>>> from mirror_gen import looking_glass
>>> with looking_glass() as what:      ❶
...         print('Alice, Kitty and Snowdrop')
...         print(what)
...
pordwonS dna yttiK ,ecilA
YKCOWREBBAJ
>>> what
'JABBERWOCKY'
>>> print('back to normal')
back to normal
```

❶ The only difference from Example 18-2 is the name of the context manager: `looking_glass` instead of `LookingGlass`.

The `contextlib.contextmanager` decorator wraps the function in a class that implements the `__enter__` and `__exit__` methods.[3]

The `__enter__` method of that class:

1. Calls the generator function to get a generator object—let's call it `gen`.
2. Calls `next(gen)` to drive it to the `yield` keyword.
3. Returns the value yielded by `next(gen)`, to allow the user to bind it to a variable in the `with/as` form.

When the `with` block terminates, the `__exit__` method:

1. Checks whether an exception was passed as `exc_type`; if so, `gen.throw(exception)` is invoked, causing the exception to be raised in the `yield` line inside the generator function body.
2. Otherwise, `next(gen)` is called, resuming the execution of the generator function body after the `yield`.

Example 18-5 has a flaw: if an exception is raised in the body of the `with` block, the Python interpreter will catch it and raise it again in the `yield` expression inside `looking_glass`. But there is no error handling there, so the `looking_glass` generator will terminate without ever restoring the original `sys.stdout.write` method, leaving the system in an invalid state.

Example 18-7 adds special handling of the `ZeroDivisionError` exception, making it functionally equivalent to the class-based Example 18-3.

Example 18-7. mirror_gen_exc.py: generator-based context manager implementing exception handling—same external behavior as Example 18-3

```
import contextlib
import sys


@contextlib.contextmanager
def looking_glass():
    original_write = sys.stdout.write

    def reverse_write(text):
```

```
            original_write(text[::-1])

        sys.stdout.write = reverse_write
        msg = ''     ❶
        try:
            yield 'JABBERWOCKY'
        except ZeroDivisionError:   ❷
            msg = 'Please DO NOT divide by zero!'
        finally:
            sys.stdout.write = original_write   ❸
            if msg:
                print(msg)    ❹
```

❶ Create a variable for a possible error message; this is the first change in relation to [Example 18-5](#).

❷ Handle `ZeroDivisionError` by setting an error message.

❸ Undo monkey-patching of `sys.stdout.write`.

❹ Display error message, if it was set.

Recall that the `__exit__` method tells the interpreter that it has handled the exception by returning a truthy value; in that case, the interpreter suppresses the exception. On the other hand, if `__exit__` does not explicitly return a value, the interpreter gets the usual `None`, and propagates the exception. With `@contextmanager`, the default behavior is inverted: the `__exit__` method provided by the decorator assumes any exception sent into the generator is handled and should be suppressed.

TIP

Having a `try/finally` (or a `with` block) around the `yield` is an unavoidable price of using `@contextmanager`, because you never know what the users of your context manager are going to do inside the `with` block.[4]

A little-known feature of `@contextmanager` is that the generators decorated with it can also be used as decorators themselves.[5] That happens because `@contextmanager` is implemented with the `contextlib.ContextDecorator` class.

Example 18-8 shows the `looking_glass` context manager from Example 18-5 used as decorator.

Example 18-8. The `looking_glass` context manager also works as a decorator

```
>>> @looking_glass()
... def verse():
...     print('The time has come')
...
>>> verse()        ❶
emoc sah emit ehT
>>> print('back to normal')        ❷
back to normal
```

❶ `looking_glass` does its job before and after the body of `verse` runs.

❷ This confirms that the original `sys.write` was restored.

Contrast Example 18-8 with Example 18-6, where `looking_glass` is used as a context manager.

An interesting real-life example of `@contextmanager` outside of the standard library is Martijn Pieters' in-place file rewriting using a context manager. Example 18-9 shows how it's used.

Example 18-9. A context manager for rewriting files in place

```
import csv

with inplace(csvfilename, 'r', newline='') as (infh, outfh):
    reader = csv.reader(infh)
    writer = csv.writer(outfh)

    for row in reader:
        row += ['new', 'columns']
        writer.writerow(row)
```

The `inplace` function is a context manager that gives you two handles — `infh` and `outfh` in the example—to the same file, allowing your code

to read and write to it at the same time. It's easier to use than the standard library's `fileinput.input` [function](#) (which also provides a context manager, by the way).

If you want to study Martijn's `inplace` source code (listed in [the post](#)), find the `yield` keyword: everything before it deals with setting up the context, which entails creating a backup file, then opening and yielding references to the readable and writable file handles that will be returned by the `__enter__` call. The `__exit__` processing after the `yield` closes the file handles and restores the file from the backup if something went wrong.

This concludes our overview of the `with` statement and context managers. Let's turn to `match/case` in the context of a complete example.

Pattern Matching in lis.py: A Case Study

In ["Pattern Matching Sequences in an Interpreter"](#) we saw examples of sequence patterns extracted from the `evaluate` function of Peter Norvig's *lis.py* interpreter, ported to Python 3.10. In this section I want to give a broader overview of how *lis.py* works, and also explore all the `case` clauses of `evaluate`, explaining not only the patterns but also what the interpreter does in each `case`.

Besides showing more pattern matching, I wrote this section for three reasons:

1. Norvig's *lis.py* is a beautiful example of idiomatic Python code.
2. The simplicity of Scheme is a master class of language design.
3. Learning how an interpreter works gave me a deeper understanding of Python and programming languages in general—interpreted or compiled.

Before looking at the Python code, let's get a little taste of Scheme so you can make sense of this case study—in case you haven't seen Scheme or Lisp before.

Scheme Syntax

In Scheme there is no distinction between expressions and statements, like we have in Python. Also, there are no infix operators. All expressions use prefix notation like `(+ x 13)` instead of `x + 13`. The same prefix notation is used for function calls—e.g., `(gcd x 13)` —and special forms —e.g., `(define x 13)`, which we'd write as the assignment statement `x = 13` in Python. The notation used by Scheme and most Lisp dialects is known as *S-expression*.[6]

Example 18-10 shows a simple example in Scheme.

Example 18-10. Greatest common divisor in Scheme

```scheme
(define (mod m n)
    (- m (* n (quotient m n)))) 

(define (gcd m n)
    (if (= n 0)
        m
        (gcd n (mod m n)))) 

(display (gcd 18 45))
```

Example 18-10 shows three Scheme expressions: two function definitions — `mod` and `gcd` —and a call to `display`, which will output 9, the result of `(gcd 18 45)`. Example 18-11 is the same code in Python (shorter than an English explanation of the recursive *Euclidean algorithm*).

Example 18-11. Same as Example 18-10, written in Python

```python
def mod(m, n):
    return m - (m // n * n)

def gcd(m, n):
    if n == 0:
        return m
    else:
        return gcd(n, mod(m, n))

print(gcd(18, 45))
```

In idiomatic Python, I'd use the `%` operator instead of reinventing `mod`, and it would be more efficient to use a `while` loop instead of recursion. But I wanted to show two function definitions, and make the examples as similar as possible, to help you read the Scheme code.

Scheme has no iterative control flow commands like `while` or `for`. Iteration is done with recursion. Note how there are no assignments in the Scheme and Python examples. Extensive use of recursion and minimal use of assignment are hallmarks of programming in a functional style.[7]

Now let's review the code of the Python 3.10 version of *lis.py*. The complete source code with tests is in the *18-with-match/lispy/py3.10/* directory of the GitHub repository *fluentpython/example-code-2e*.

Imports and Types

Example 18-12 shows the first lines of *lis.py*. The use of `TypeAlias` and the `|` type union operator require Python 3.10.

Example 18-12. lis.py: top of the file

```python
import math
import operator as op
from collections import ChainMap
from itertools import chain
from typing import Any, TypeAlias, NoReturn

Symbol: TypeAlias = str
Atom: TypeAlias = float | int | Symbol
Expression: TypeAlias = Atom | list
```

The types defined are:

Symbol

> Just an alias for `str`. In *lis.py*, `Symbol` is used for identifiers; there is no string data type with operations such as slicing, splitting, etc.[8]

Atom

A simple syntactic element, such as a number or a `Symbol` —as opposed to a composite structure made of distinct parts, like a list.

Expression

The building blocks of Scheme programs are expressions made of atoms and lists, possibly nested.

The Parser

Norvig's parser is 36 lines of code showcasing the power of Python applied to handling the simple recursive syntax of S-expression—without string data, comments, macros, and other features of standard Scheme that make parsing more complicated (Example 18-13).

Example 18-13. lis.py: the main parsing functions

```python
def parse(program: str) -> Expression:
    "Read a Scheme expression from a string."
    return read_from_tokens(tokenize(program))

def tokenize(s: str) -> list[str]:
    "Convert a string into a list of tokens."
    return s.replace('(', ' ( ').replace(')', ' ) ').split()

def read_from_tokens(tokens: list[str]) -> Expression:
    "Read an expression from a sequence of tokens."
    # more parsing code omitted in book listing
```

The main function of that group is `parse`, which takes an S-expression as a `str` and returns an `Expression` object, as defined in Example 18-12: an `Atom` or a `list` that may contain more atoms and nested lists.

Norvig uses a smart trick in `tokenize`: he adds spaces before and after each parenthesis in the input and then splits it, resulting in a list of syntactic tokens with `'('` and `')'` as separate tokens. This shortcut works because there is no string type in the little Scheme of *lis.py*, so every `'('` or `')'` is an expression delimiter. The recursive parsing code is in `read_from_tokens`, a 14-line function that you can read in the *fluentpython/example-code-2e* repository. I will skip it because I want to focus on the other parts of the interpreter.

Here are some doctests extracted from *lispy/py3.10/examples_test.py*:

```
>>> from lis import parse
>>> parse('1.5')
1.5
>>> parse('ni!')
'ni!'
>>> parse('(gcd 18 45)')
['gcd', 18, 45]
>>> parse('''
... (define double
...       (lambda (n)
...            (* n 2)))
... ''')
['define', 'double', ['lambda', ['n'], ['*', 'n', 2]]]
```

The parsing rules for this subset of Scheme are simple:

1. A token that looks like a number is parsed as a `float` or `int`.
2. Anything else that is not `'('` or `')'` is parsed as a `Symbol` —a `str` to be used as an identifier. This includes source text like `+`, `set!`, and `make-counter` that are valid identifiers in Scheme but not in Python.
3. Expressions inside `'('` and `')'` are recursively parsed as lists containing atoms or as nested lists that may contain atoms and more nested lists.

Using the terminology of the Python interpreter, the output of `parse` is an AST (Abstract Syntax Tree): a convenient representation of the Scheme program as nested lists forming a tree-like structure, where the outermost list is the trunk, inner lists are the branches, and atoms are the leaves ([Figure 18-1](#)).

Figure 18-1. A Scheme `lambda` expression represented as source code (concrete syntax), as a tree, and as a sequence of Python objects (abstract syntax).

The Environment

The `Environment` class extends `collections.ChainMap`, adding a `change` method to update a value inside one of the chained dicts, which

`ChainMap` instances hold in a list of mappings: the `self.maps` attribute. The `change` method is needed to support the Scheme `(set! …)` form, described later; see Example 18-14.

Example 18-14. *lis.py*: the `Environment` **class**

```python
class Environment(ChainMap[Symbol, Any]):
    "A ChainMap that allows changing an item in-place."

    def change(self, key: Symbol, value: Any) -> None:
        "Find where key is defined and change the value there."
        for map in self.maps:
            if key in map:
                map[key] = value  # type: ignore[index]
                return
        raise KeyError(key)
```

Note that the `change` method only updates existing keys.[9] Trying to change a key that is not found raises `KeyError`.

This doctest shows how `Environment` works:

```python
>>> from lis import Environment
>>> inner_env = {'a': 2}
>>> outer_env = {'a': 0, 'b': 1}
>>> env = Environment(inner_env, outer_env)
>>> env['a']       ❶
2
>>> env['a'] = 111   ❷
>>> env['c'] = 222
>>> env
Environment({'a': 111, 'c': 222}, {'a': 0, 'b': 1})
>>> env.change('b', 333)   ❸
>>> env
Environment({'a': 111, 'c': 222}, {'a': 0, 'b': 333})
```

❶ When reading values, `Environment` works as `ChainMap`: keys are searched in the nested mappings from left to right. That's why the value of `a` in the `outer_env` is shadowed by the value in `inner_env`.

❷ Assigning with `[]` overwrites or inserts new items, but always in the first mapping, `inner_env` in this example.

❸ `env.change('b', 333)` seeks the `'b'` key and assigns a new value to it in-place, in the `outer_env`.

Next is the `standard_env()` function, which builds and returns an `Environment` loaded with predefined functions, similar to Python's `__builtins__` module that is always available (Example 18-15).

Example 18-15. lis.py: `standard_env()` **builds and returns the global environment**

```python
def standard_env() -> Environment:
    "An environment with some Scheme standard procedures."
    env = Environment()
    env.update(vars(math))    # sin, cos, sqrt, pi, ...
    env.update({
            '+': op.add,
            '-': op.sub,
            '*': op.mul,
            '/': op.truediv,
            # omitted here: more operator definitions
            'abs': abs,
            'append': lambda *args: list(chain(*args)),
            'apply': lambda proc, args: proc(*args),
            'begin': lambda *x: x[-1],
            'car': lambda x: x[0],
            'cdr': lambda x: x[1:],
            # omitted here: more function definitions
            'number?': lambda x: isinstance(x, (int, float)),
            'procedure?': callable,
            'round': round,
            'symbol?': lambda x: isinstance(x, Symbol),
    })
    return env
```

To summarize, the `env` mapping is loaded with:

- All functions from Python's `math` module
- Selected operators from Python's `op` module
- Simple but powerful functions built with Python's `lambda`

- Python built-ins renamed, like `callable` as `procedure?`, or directly mapped, like `round`

The REPL

Norvig's REPL (read-eval-print-loop) is easy to understand but not user-friendly (see Example 18-16). If no command-line arguments are given to *lis.py*, the `repl()` function is invoked by `main()` —defined at the end of the module. At the `lis.py>` prompt, we must enter correct and complete expressions; if we forget to close one parenthesis, *lis.py* crashes.[10]

Example 18-16. The REPL functions

```python
def repl(prompt: str = 'lis.py> ') -> NoReturn:
    "A prompt-read-eval-print loop."
    global_env = Environment({}, standard_env())
    while True:
        ast = parse(input(prompt))
        val = evaluate(ast, global_env)
        if val is not None:
            print(lispstr(val))

def lispstr(exp: object) -> str:
    "Convert a Python object back into a Lisp-readable string."
    if isinstance(exp, list):
        return '(' + ' '.join(map(lispstr, exp)) + ')'
    else:
        return str(exp)
```

Here is a quick explanation about these two functions:

repl(prompt: str = 'lis.py> ') -> NoReturn

Calls `standard_env()` to provide built-in functions for the global environment, then enters an infinite loop, reading and parsing each input line, evaluating it in the global environment, and displaying the result—unless it's `None`. The `global_env` may be modified by `evaluate`. For example, when a user defines a new global variable or named function, it is stored in the first mapping of the environment—the empty `dict` in the `Environment` constructor call in the first line of `repl`.

```
lispstr(exp: object) -> str
```

> The inverse function of `parse` : given a Python object representing an expression, `parse` returns the Scheme source code for it. For example, given `['+', 2, 3]`, the result is `'(+ 2 3)'`.

The Evaluator

Now we can appreciate the beauty of Norvig's expression evaluator—made a little prettier with `match/case`. The `evaluate` function in Example 18-17 takes an `Expression` built by `parse` and an `Environment`.

The body of `evaluate` is a single `match` statement with an expression `exp` as the subject. The `case` patterns express the syntax and semantics of Scheme with amazing clarity.

Example 18-17. `evaluate` **takes an expression and computes its value**

```python
KEYWORDS = ['quote', 'if', 'lambda', 'define', 'set!']

def evaluate(exp: Expression, env: Environment) -> Any:
    "Evaluate an expression in an environment."
    match exp:
        case int(x) | float(x):
            return x
        case Symbol(var):
            return env[var]
        case ['quote', x]:
            return x
        case ['if', test, consequence, alternative]:
            if evaluate(test, env):
                return evaluate(consequence, env)
            else:
                return evaluate(alternative, env)
        case ['lambda', [*parms], *body] if body:
            return Procedure(parms, body, env)
        case ['define', Symbol(name), value_exp]:
            env[name] = evaluate(value_exp, env)
        case ['define', [Symbol(name), *parms], *body] if body:
            env[name] = Procedure(parms, body, env)
        case ['set!', Symbol(name), value_exp]:
```

```
            env.change(name, evaluate(value_exp, env))
        case [func_exp, *args] if func_exp not in KEYWORDS:
            proc = evaluate(func_exp, env)
            values = [evaluate(arg, env) for arg in args]
            return proc(*values)
        case _:
            raise SyntaxError(lispstr(exp))
```

Let's study each `case` clause and what it does. In some cases I added comments showing an S-expression that would match the pattern when parsed into a Python list. Doctests extracted from *examples_test.py* demonstrate each `case`.

Evaluating numbers

```
case int(x) | float(x):
    return x
```

Subject:

Instance of `int` or `float`.

Action:

Return value as is.

Example:

```
>>> from lis import parse, evaluate, standard_env
>>> evaluate(parse('1.5'), {})
1.5
```

Evaluating symbols

```
case Symbol(var):
    return env[var]
```

Subject:

Instance of `Symbol`, i.e., a `str` used as an identifier.

Look up `var` in `env` and return its value.

Examples:

```
>>> evaluate(parse('+'), standard_env())
<built-in function add>
>>> evaluate(parse('ni!'), standard_env())
Traceback (most recent call last):
    ...
KeyError: 'ni!'
```

(quote ...)

The `quote` special form treats atoms and lists as data instead of expressions to be evaluated.

```
# (quote (99 bottles of beer))
case ['quote', x]:
    return x
```

Subject:

List starting with the symbol `'quote'`, followed by one expression `x`.

Action:

Return `x` without evaluating it.

Examples:

```
>>> evaluate(parse('(quote no-such-name)'), standard_env())
'no-such-name'
>>> evaluate(parse('(quote (99 bottles of beer))'), standard_env())
[99, 'bottles', 'of', 'beer']
>>> evaluate(parse('(quote (/ 10 0))'), standard_env())
['/', 10, 0]
```

Without `quote`, each expression in the test would raise an error:

- `no-such-name` would be looked up in the environment, raising `KeyError`
- `(99 bottles of beer)` cannot be evaluated because the number 99 is not a `Symbol` naming a special form, operator, or function
- `(/ 10 0)` would raise `ZeroDivisionError`

WHY LANGUAGES HAVE RESERVED KEYWORDS

Although simple, `quote` cannot be implemented as a function in Scheme. Its special power is to prevent the interpreter from evaluating `(f 10)` in the expression `(quote (f 10))`: the result is simply a list with a `Symbol` and an `int`. In contrast, in a function call like `(abs (f 10))`, the interpreter evaluates `(f 10)` before invoking `abs`. That's why `quote` is a reserved keyword: it must be handled as a special form.

In general, reserved keywords are needed:

- To introduce specialized evaluation rules, as in `quote` and `lambda` —which don't evaluate any of their subexpressions
- To change the control flow, as in `if` and function calls—which also have special evaluation rules
- To manage the environment, as in `define` and `set`

This is also why Python, and programming languages in general, need reserved keywords. Think about Python's `def`, `if`, `yield`, `import`, `del`, and what they do.

(if …)

```python
# (if (< x 0) 0 x)
case ['if', test, consequence, alternative]:
    if evaluate(test, env):
        return evaluate(consequence, env)
    else:
        return evaluate(alternative, env)
```

Subject:

List starting with `'if'` followed by three expressions: `test`, `consequence`, and `alternative`.

Action:

Evaluate `test`:

- If true, evaluate `consequence` and return its value.
- Otherwise, evaluate `alternative` and return its value.

Examples:

```
>>> evaluate(parse('(if (= 3 3) 1 0))'), standard_env())
1
>>> evaluate(parse('(if (= 3 4) 1 0))'), standard_env())
0
```

The `consequence` and `alternative` branches must be single expressions. If more than one expression is needed in a branch, you can combine them with `(begin exp1 exp2…)`, provided as a function in *lis.py*—see [Example 18-15](#).

(lambda ...)

Scheme's `lambda` form defines anonymous functions. It doesn't suffer from the limitations of Python's `lambda`: any function that can be written in Scheme can be written using the `(lambda …)` syntax.

```
# (lambda (a b) (/ (+ a b) 2))
case ['lambda' [*parms], *body] if body:
    return Procedure(parms, body, env)
```

Subject:

List starting with `'lambda'`, followed by:

- List of zero or more parameter names.
- One or more expressions collected in `body` (the guard ensures that `body` is not empty).

Action:

Create and return a new `Procedure` instance with the parameter names, the list of expressions as the body, and the current environment.

Example:

```
>>> expr = '(lambda (a b) (* (/ a b) 100))'
>>> f = evaluate(parse(expr), standard_env())
>>> f   # doctest: +ELLIPSIS
<lis.Procedure object at 0x...>
>>> f(15, 20)
75.0
```

The `Procedure` class implements the concept of a closure: a callable object holding parameter names, a function body, and a reference to the environment in which the function is defined. We'll study the code for `Procedure` in a moment.

(define …)

The `define` keyword is used in two different syntactic forms. The simplest is:

```
# (define half (/ 1 2))
case ['define', Symbol(name), value_exp]:
    env[name] = evaluate(value_exp, env)
```

Subject:

List starting with `'define'`, followed by a `Symbol` and an expression.

Action:

Evaluate the expression and put its value into `env`, using `name` as key.

Example:

```
>>> global_env = standard_env()
>>> evaluate(parse('(define answer (* 7 6))'), global_env)
```

```
>>> global_env['answer']
42
```

The doctest for this `case` creates a `global_env` so that we can verify that `evaluate` puts `answer` into that `Environment`.

We can use that simple `define` form to create variables or to bind names to anonymous functions, using `(lambda …)` as the `value_exp`.

Standard Scheme provides a shortcut for defining named functions. That's the second `define` form:

```
# (define (average a b) (/ (+ a b) 2))
case ['define', [Symbol(name), *parms], *body] if body:
    env[name] = Procedure(parms, body, env)
```

Subject:

List starting with `'define'`, followed by:

- A list starting with a `Symbol(name)`, followed by zero or more items collected into a list named `parms`.
- One or more expressions collected in `body` (the guard ensures that `body` is not empty).

Action:

- Create a new `Procedure` instance with the parameter names, the list of expressions as the body, and the current environment.
- Put the `Procedure` into `env`, using `name` as key.

The doctest in Example 18-18 defines a function named `%` that computes a percentage and adds it to the `global_env`.

Example 18-18. Defining a function named `%` that computes a percentage

```
>>> global_env = standard_env()
>>> percent = '(define (% a b) (* (/ a b) 100))'
>>> evaluate(parse(percent), global_env)
```

```
>>> global_env['%']  # doctest: +ELLIPSIS
<lis.Procedure object at 0x...>
>>> global_env['%'](170, 200)
85.0
```

After calling `evaluate`, we check that `%` is bound to a `Procedure` that takes two numeric arguments and returns a percentage.

The pattern for the second `define` case does not enforce that the items in `parms` are all `Symbol` instances. I'd have to check that before building the `Procedure`, but I didn't—to keep the code as easy to follow as Norvig's.

(set! …)

The `set!` form changes the value of a previously defined variable.[11]

```
    # (set! n (+ n 1))
    case ['set!', Symbol(name), value_exp]:
        env.change(name, evaluate(value_exp, env))
```

Subject:

List starting with `'set!'`, followed by a `Symbol` and an expression.

Action:

Update the value of `name` in `env` with the result of evaluating the expression.

The `Environment.change` method traverses the chained environments from local to global, and updates the first occurrence of `name` with the new value. If we were not implementing the `'set!'` keyword, we could use Python's `ChainMap` as the `Environment` type everywhere in this interpreter.

The use of the `set!` form is related to the use of the `nonlocal` keyword in Python: declaring `nonlocal x` allows `x = 10` to update a previously

defined `x` variable outside of the local scope. Without a `nonlocal x` declaration, `x = 10` will always create a local variable in Python, as we saw in ["The nonlocal Declaration"](#).

Similarly, `(set! x 10)` updates a previously defined `x` that may be outside of the local environment of the function. In contrast, the variable `x` in `(define x 10)` is always a local variable, created or updated in the local environment.

Both `nonlocal` and `(set! …)` are needed to update program state held in variables within a closure. [Example 9-13](#) demonstrated the use of `nonlocal` to implement a function to compute a running average, holding an item `count` and `total` in a closure. Here is that same idea, written in the Scheme subset of *lis.py*:

```
(define (make-averager)
    (define count 0)
    (define total 0)
    (lambda (new-value)
        (set! count (+ count 1))
        (set! total (+ total new-value))
        (/ total count)
    )
)
(define avg (make-averager))    ❶
(avg 10)    ❷
(avg 11)    ❸
(avg 15)    ❹
```

❶ Creates a new closure with the inner function defined by `lambda`, and the variables `count` and `total` initialized to 0; binds the closure to `avg`.

❷ Returns 10.0.

❸ Returns 10.5.

❹ Returns 12.0.

The preceding code is one of the tests in *[lispy/py3.10/examples_test.py](#)*.

Now we get to a function call.

Function call

```
# (gcd (* 2 105) 84)
case [func_exp, *args] if func_exp not in KEYWORDS:
    proc = evaluate(func_exp, env)
    values = [evaluate(arg, env) for arg in args]
    return proc(*values)
```

Subject:

List with one or more items.

The guard ensures that `func_exp` is not one of `['quote', 'if', 'define', 'lambda', 'set!']` —listed right before `evaluate` in Example 18-17.

The pattern matches any list with one or more expressions, binding the first expression to `func_exp` and the rest to `args` as a list, which may be empty.

Action:

- Evaluate `func_exp` to obtain a function `proc`.
- Evaluate each item in `args` to build a list of argument values.
- Call `proc` with the values as separate arguments, returning the result.

Example:

```
>>> evaluate(parse('(% (* 12 14) (- 500 100))'), global_env)
42.0
```

This doctest continues from Example 18-18: it assumes `global_env` has a function named `%`. The arguments given to `%` are arithmetic expressions, to emphasize that the arguments are evaluated before the function is called.

The guard in this `case` is needed because `[func_exp, *args]`
matches any sequence subject with one or more items. However, if
`func_exp` is a keyword, and the subject did not match any previous
case, then it is really a syntax error.

Catch syntax errors

If the subject `exp` does not match any of the previous cases, the catch-all
`case` raises a `SyntaxError`:

```
    case _:
        raise SyntaxError(lispstr(exp))
```

Here is an example of a malformed `(lambda …)` reported as a
`SyntaxError`:

```
  >>> evaluate(parse('(lambda is not like this)'), standard_env())
  Traceback (most recent call last):
    ...
  SyntaxError: (lambda is not like this)
```

If the `case` for function call did not have that guard rejecting keywords,
the `(lambda is not like this)` expression would be handled as a
function call, which would raise `KeyError` because `'lambda'` is not
part of the environment—just like `lambda` is not a Python built-in func-
tion.

Procedure: A Class Implementing a Closure

The `Procedure` class could very well be named `Closure`, because that's
what it represents: a function definition together with an environment.
The function definition includes the name of the parameters and the ex-
pressions that make up the body of the function. The environment is used
when the function is called to provide the values of the *free variables*:
variables that appear in the body of the function but are not parameters,
local variables, or global variables. We saw the concepts of *closure* and
free variable in ["Closures"](#).

We learned how to use closures in Python, but now we can dive deeper and see how a closure is implemented in *lis.py*:

```python
class Procedure:
    "A user-defined Scheme procedure."

    def __init__(  ❶
        self, parms: list[Symbol], body: list[Expression], env: Environme
    ):
        self.parms = parms   ❷
        self.body = body
        self.env = env

    def __call__(self, *args: Expression) -> Any:   ❸
        local_env = dict(zip(self.parms, args))   ❹
        env = Environment(local_env, self.env)   ❺
        for exp in self.body:   ❻
            result = evaluate(exp, env)
        return result   ❼
```

❶ Called when a function is defined by the `lambda` or `define` forms.

❷ Save the parameter names, body expressions, and environment for later use.

❸ Called by `proc(*values)` in the last line of the `case [func_exp, *args]` clause.

❹ Build `local_env` mapping `self.parms` as local variable names, and the given `args` as values.

❺ Build a new combined `env`, putting `local_env` first, and then `self.env`—the environment that was saved when the function was defined.

❻ Iterate over each expression in `self.body`, evaluating it in the combined `env`.

❼ Return the result of the last expression evaluated.

There are a couple of simple functions after `evaluate` in *lis.py*: `run` reads a complete Scheme program and executes it, and `main` calls `run` or `repl`, depending on the command line—similar to what Python does. I will not describe those functions because there's nothing new in them. My goals were to share with you the beauty of Norvig's little interpreter, to give more insight into how closures work, and to show how `match/case` is a great addition to Python.

To wrap up this extended section on pattern matching, let's formalize the concept of an OR-pattern.

Using OR-patterns

A series of patterns separated by `|` is an *OR-pattern*: it succeeds if any of the subpatterns succeed. The pattern in "Evaluating numbers" is an OR-pattern:

```python
case int(x) | float(x):
    return x
```

All subpatterns in an OR-pattern must use the same variables. This restriction is necessary to ensure that the variables are available to the guard expression and the `case` body, regardless of the subpattern that matched.

WARNING

In the context of a `case` clause, the `|` operator has a special meaning. It does not trigger the `__or__` special method, which handles expressions like `a | b` in other contexts, where it is overloaded to perform operations such as set union or integer bitwise-or, depending on the operands.

An OR-pattern is not restricted to appear at the top level of a pattern. You can also use `|` in subpatterns. For example, if we wanted *lis.py* to accept the Greek letter λ (lambda)[12] as well as the `lambda` keyword, we can rewrite the pattern like this:

```
    # (λ (a b) (/ (+ a b) 2) )
    case ['lambda' | 'λ', [*parms], *body] if body:
        return Procedure(parms, body, env)
```

Now we can move to the third and last subject of this chapter: the un-usual places where an `else` clause may appear in Python.

Do This, Then That: else Blocks Beyond if

This is no secret, but it is an underappreciated language feature: the `else` clause can be used not only in `if` statements but also in `for`, `while`, and `try` statements.

The semantics of `for/else`, `while/else`, and `try/else` are closely related, but very different from `if/else`. Initially, the word `else` actu-ally hindered my understanding of these features, but eventually I got used to it.

Here are the rules:

for

> The `else` block will run only if and when the `for` loop runs to completion (i.e., not if the `for` is aborted with a `break`).

while

> The `else` block will run only if and when the `while` loop exits because the condition became *falsy* (i.e., not if the `while` is aborted with a `break`).

try

> The `else` block will run only if no exception is raised in the `try` block. The [official docs](#) also state: "Exceptions in the `else` clause are not handled by the preceding `except` clauses."

In all cases, the `else` clause is also skipped if an exception or a `return`, `break`, or `continue` statement causes control to jump out of the main

block of the compound statement.

Using `else` with these statements often makes the code easier to read and saves the trouble of setting up control flags or coding extra `if` statements.

The use of `else` in loops generally follows the pattern of this snippet:

```python
for item in my_list:
    if item.flavor == 'banana':
        break
else:
    raise ValueError('No banana flavor found!')
```

In the case of `try/except` blocks, `else` may seem redundant at first. After all, the `after_call()` in the following snippet will run only if the `dangerous_call()` does not raise an exception, correct?

```python
try:
    dangerous_call()
    after_call()
except OSError:
    log('OSError...')
```

However, doing so puts the `after_call()` inside the `try` block for no good reason. For clarity and correctness, the body of a `try` block should only have the statements that may generate the expected exceptions. This is better:

```python
try:
    dangerous_call()
except OSError:
    log('OSError...')
else:
    after_call()
```

Now it's clear that the `try` block is guarding against possible errors in `dangerous_call()` and not in `after_call()`. It's also explicit that `after_call()` will only execute if no exceptions are raised in the `try` block.

In Python, `try/except` is commonly used for control flow, and not just for error handling. There's even an acronym/slogan for that documented in the official Python glossary:

> EAFP
>
> > Easier to ask for forgiveness than permission. This common Python coding style assumes the existence of valid keys or attributes and catches exceptions if the assumption proves false. This clean and fast style is characterized by the presence of many try and except statements. The technique contrasts with the LBYL style common to many other languages such as C.

The glossary then defines LBYL:

> LBYL
>
> > Look before you leap. This coding style explicitly tests for pre-conditions before making calls or lookups. This style contrasts with the EAFP approach and is characterized by the presence of many if statements. In a multi-threaded environment, the LBYL approach can risk introducing a race condition between "the looking" and "the leaping." For example, the code, if key in mapping: return mapping[key] can fail if another thread removes key from mapping after the test, but before the lookup. This issue can be solved with locks or by using the EAFP approach.

Given the EAFP style, it makes even more sense to know and use `else` blocks well in `try/except` statements.

NOTE

When the `match` statement was discussed, some people (including me) thought it should also have an `else` clause. In the end it was decided that it wasn't needed because `case _:` does the same job.[13]

Now let's summarize the chapter.

Chapter Summary

This chapter started with context managers and the meaning of the `with` statement, quickly moving beyond its common use to automatically close opened files. We implemented a custom context manager: the `LookingGlass` class with the `__enter__`/`__exit__` methods, and saw how to handle exceptions in the `__exit__` method. A key point that Raymond Hettinger made in his PyCon US 2013 keynote is that `with` is not just for resource management; it's a tool for factoring out common setup and teardown code, or any pair of operations that need to be done before and after another procedure.[14]

We reviewed functions in the `contextlib` standard library module. One of them, the `@contextmanager` decorator, makes it possible to implement a context manager using a simple generator with one `yield`—a leaner solution than coding a class with at least two methods. We reimplemented the `LookingGlass` as a `looking_glass` generator function, and discussed how to do exception handling when using `@contextmanager`.

Then we studied Peter Norvig's elegant *lis.py*, a Scheme interpreter written in idiomatic Python, refactored to use `match/case` in `evaluate`—the function at the core of any interpreter. Understanding how `evaluate` works required reviewing a little bit of Scheme, a parser for S-expressions, a simple REPL, and the construction of nested scopes through an `Environment` subclass of `collection.ChainMap`. In the

end, *lis.py* became a vehicle to explore much more than pattern matching. It shows how the different parts of an interpreter work together, illuminating core features of Python itself: why reserved keywords are necessary, how scoping rules work, and how closures are built and used.

Further Reading

[Chapter 8, "Compound Statements,"](#) in *The Python Language Reference* says pretty much everything there is to say about `else` clauses in `if`, `for`, `while`, and `try` statements. Regarding Pythonic usage of `try/except`, with or without `else`, Raymond Hettinger has a brilliant answer to the question ["Is it a good practice to use try-except-else in Python?"](#) in StackOverflow. *[Python in a Nutshell](#)*, 3rd ed., by Martelli et al., has a chapter about exceptions with an excellent discussion of the EAFP style, crediting computing pioneer Grace Hopper for coining the phrase, "It's easier to ask forgiveness than permission."

The Python Standard Library, Chapter 4, "Built-in Types," has a section devoted to ["Context Manager Types"](#). The `__enter__`/`__exit__` special methods are also documented in *The Python Language Reference* in ["With Statement Context Managers"](#). Context managers were introduced in [PEP 343—The "with" Statement](#).

Raymond Hettinger highlighted the `with` statement as a "winning language feature" in his [PyCon US 2013 keynote](#). He also showed some interesting applications of context managers in his talk, ["Transforming Code into Beautiful, Idiomatic Python"](#), at the same conference.

Jeff Preshing's blog post ["The Python *with* Statement by Example"](#) is interesting for the examples using context managers with the `pycairo` graphics library.

The `contextlib.ExitStack` class is based on an original idea by Nikolaus Rath, who wrote a short post explaining why its useful: ["On the Beauty of Python's ExitStack"](#). In that text, Rath submits that `ExitStack` is similar but more flexible than the `defer` statement in Go—which I think is one of the best ideas in that language.

Beazley and Jones devised context managers for very different purposes in their *Python Cookbook,* 3rd ed. "Recipe 8.3. Making Objects Support the Context-Management Protocol" implements a `LazyConnection` class whose instances are context managers that open and close network connections automatically in `with` blocks. "Recipe 9.22. Defining Context Managers the Easy Way" introduces a context manager for timing code, and another for making transactional changes to a `list` object: within the `with` block, a working copy of the `list` instance is made, and all changes are applied to that working copy. Only when the `with` block completes without an exception, the working copy replaces the original list. Simple and ingenious.

Peter Norvig describes his small Scheme interpreters in the posts "(How to Write a (Lisp) Interpreter (in Python))" and "(An ((Even Better) Lisp) Interpreter (in Python))". The code for *lis.py* and *lispy.py* is the *norvig/pytudes* repository. My repository *fluentpython/lispy* includes the *mylis* forks of *lis.py*, updated to Python 3.10, with a nicer REPL, command-line integration, examples, more tests, and references for learning more about Scheme. The best Scheme dialect and environment to learn and experiment is Racket.

Factoring Out the Bread

In his PyCon US 2013 keynote, ["What Makes Python Awesome"](), Raymond Hettinger says when he first saw the `with` statement proposal he thought it was "a little bit arcane." Initially, I had a similar reaction. PEPs are often hard to read, and PEP 343 is typical in that regard.

Then—Hettinger told us—he had an insight: subroutines are the most important invention in the history of computer languages. If you have sequences of operations like A;B;C and P;B;Q, you can factor out B in a subroutine. It's like factoring out the filling in a sandwich: using tuna with different breads. But what if you want to factor out the bread, to make sandwiches with wheat bread, using a different filling each time? That's what the `with` statement offers. It's the complement of the subroutine. Hettinger went on to say:

> *The* `with` *statement is a very big deal. I encourage you to go out and take this tip of the iceberg and drill deeper. You can probably do profound things with the* `with` *statement. The best uses of it have not been discovered yet. I expect that if you make good use of it, it will be copied into other languages and all future languages will have it. You can be part of discovering something almost as profound as the invention of the subroutine itself.*

Hettinger admits he is overselling the `with` statement. Nevertheless, it is a very useful feature. When he used the sandwich analogy to explain how `with` is the complement to the subroutine, many possibilities opened up in my mind.

If you need to convince anyone that Python is awesome, you should watch Hettinger's keynote. The bit about context managers is from 23:00 to 26:15. But the entire keynote is excellent.

Efficient Recursion with Proper Tail Calls

Standard Scheme implementations are required to provide *proper tail calls* (PTC), to make iteration through recursion a practical alternative to `while` loops in imperative languages. Some writers refer to PTC as *tail call optimization* (TCO); for others, TCO is something different. For more details, see ["Tail call"](#) on Wikipedia and ["Tail call optimization in ECMAScript 6"](#).

A *tail call* is when a function returns the result of a function call, which may be the same function or not. The `gcd` examples in [Example 18-10](#) and [Example 18-11](#) make (recursive) tail calls in the *falsy* branch of the `if`.

On the other hand, this `factorial` does not make a tail call:

```python
def factorial(n):
    if n < 2:
        return 1
    return n * factorial(n - 1)
```

The call to `factorial` in the last line is not a tail call because the `return` value is not the result of the recursive call: the result is multi-

plied by `n` before it is returned.

Here is an alternative that uses a tail call, and is therefore *tail recursive*:

```python
def factorial_tc(n, product=1):
    if n < 1:
        return product
    return factorial_tc(n - 1, product * n)
```

Python does not have PTC, so there's no advantage in writing tail recursive functions. In this case, the first version is shorter and more readable in my opinion. For real-life uses, don't forget that Python has `math.factorial`, written in C without recursion. The point is that, even in languages that implement PTC, it does not benefit every recursive function, only those that are carefully written to make tail calls.

If PTC is supported by the language, when the interpreter sees a tail call, it jumps into the body of the called function without creating a new stack frame, saving memory. There are also compiled languages that implement PTC, sometimes as an optimization that can be toggled.

There is no universal consensus about the definition of TCO or the value of PTC in languages that were not designed as functional languages from the ground up, like Python or JavaScript. In functional languages, PTC is an expected feature, not merely an optimization that is nice to have. If a language has no iteration mechanism other than recursion, then PTC is necessary for practical usage. Norvig's *lis.py* does not implement PTC, but his more elaborate *lispy.py* interpreter does.

The Case Against Proper Tail Calls in Python and JavaScript

CPython does not implement PTC, and probably never will. Guido van Rossum wrote "Final Words on Tail Calls" to explain why. To summarize, here is a key passage from his post:

> *Personally, I think it is a fine feature for some languages, but I don't think it fits Python: the elimination of stack traces for some calls but not others would certainly confuse many users, who have not been raised with tail call religion but might have learned about call semantics by tracing through a few calls in a debugger.*

In 2015, PTC was included in the ECMAScript 6 standard for JavaScript. As of October 2021, the interpreter in [WebKit implements it](#). WebKit is used by Safari. The JS interpreters in every other major browser don't have PTC, and neither does Node.js, as it relies on the V8 engine that Google maintains for Chrome. Transpilers and polyfills targeting JS, like TypeScript, ClojureScript, and Babel, don't support PTC either, according to this ["ECMAScript 6 compatibility table"](#).

I've seen several explanations for the rejection of PTC by the implementers, but the most common is the same that Guido van Rossum mentioned: PTC makes debugging harder for everyone, while benefiting only a minority of people who'd rather use recursion for iteration. For details, see ["What happened to proper tail calls in JavaScript?"](#) by Graham Marlow.

There are cases when recursion is the best solution, even in Python without PTC. In a [previous post](#) on the subject, Guido wrote:

> [...] a typical Python implementation allows 1000 recursions, which is plenty for non-recursively written code and for code that recourses to traverse, for example, a typical parse tree, but not enough for a recursively written loop over a large list.

I agree with Guido and the majority of JS implementers: PTC is not a good fit for Python or JavaScript. The lack of PTC is the main restriction for writing Python programs in a functional style—more than the limited `lambda` syntax.

If you are curious to see how PTC works in an interpreter with less features (and less code) than Norvig's *lispy.py*, check out *[mylis_2](#)*. The trick starts with the infinite loop in `evaluate` and the code in the `case` for function calls: that combination makes the intepreter jump into the body of the next `Procedure` without calling `evaluate` recursively during a tail call. Those little interpreters demonstrate the power of abstraction: even though Python does not implement PTC, it's possible and not very hard to write an interpreter, in Python, that does implement PTC. I learned how to do it reading Peter Norvig's code. Thanks for sharing it, professor!

Norvig's Take on evaluate() with Pattern Matching

I shared the code for the Python 3.10 version of *lis.py* with Peter Norvig. He liked the example using pattern matching, but suggested a different solution: instead of the guards I wrote, he would have exactly one `case` per keyword, and have tests within each `case`, to provide more specific `SyntaxError` messages—for example, when a body is empty. This would also make the guard in `case [func_exp, *args] if func_exp not in KEYWORDS:` unnecessary, as every keyword would be handled before the `case` for function calls.

I'll probably follow Norvig's advice when I add more functionality to *mylis*. But the way I structured `evaluate` in Example 18-17 has some didactic advantages for this book: the example parallels the implementation with `if/elif/…` (Example 2-11), the `case` clauses demonstrate more features of pattern matching, and the code is more concise.

---

1   PyCon US 2013 keynote: "What Makes Python Awesome"; the part about `with` starts at 23:00 and ends at 26:15.

2   The three arguments received by `self` are exactly what you get if you call `sys.exc_info()` in the `finally` block of a `try/finally` statement. This makes sense, considering that the `with` statement is meant to replace most uses of `try/finally`, and calling `sys.exc_info()` was often necessary to determine what clean-up action would be required.

3   The actual class is named `_GeneratorContextManager`. If you want to see exactly how it works, read its source code in *Lib/contextlib.py* in Python 3.10.

4   This tip is quoted literally from a comment by Leonardo Rochael, one of the tech reviewers for this book. Nicely said, Leo!

5   At least I and the other technical reviewers didn't know it until Caleb Hattingh told us. Thanks, Caleb!

6   People complain about too many parentheses in Lisp, but thoughtful indentation and a good editor mostly take care of that issue. The main readability problem is using the same `(f …)` notation for function calls and special forms like `(define …)`, `(if …)`, and `(quote …)` that don't behave at all like function calls.

7   To make iteration through recursion practical and efficient, Scheme and other functional languages implement *proper tail calls*. For more about this, see "Soapbox".

8   But Norvig's second interpreter, *lispy.py*, supports strings as a data type, as well as advanced features like syntactic macros, continuations, and proper tail calls. However, *lispy.py* is almost three times longer than *lis.py*—and much harder to understand.

9   The `# type: ignore[index]` comment is there because of *typeshed* issue #6042, which is unresolved as I review this chapter. `ChainMap` is annotated as `MutableMapping`, but the type hint in the `maps` attribute says it's a list of `Mapping`, indirectly making the whole `ChainMap` immutable as far as Mypy is concerned.

10   As I studied Norvig's *lis.py* and *lispy.py*, I started a fork named *mylis* that adds some features, including a REPL that accepts partial S-expressions and prompts for the continuation, similar to how Python's REPL knows we are not finished and presents the secondary prompt (. . .) until we enter a complete expression or statement that can be evaluated. *mylis* also handles a few errors gracefully, but it's still easy to crash. It's not nearly as robust as Python's REPL.

11   Assignment is one of the first features taught in many programming tutorials, but `set!` only appears on page 220 of the best known Scheme book, *Structure and Interpretation of Computer Programs,* 2nd ed., by Abelson et al. (MIT Press), a.k.a. SICP or the "Wizard Book." Coding in a functional style can take us very far without the state changes that are typical of imperative and object-oriented programming.

12   The official Unicode name for λ (U+03BB) is GREEK SMALL LETTER LAMDA. This is not a typo: the character is named "lamda" without the "b" in the Unicode database. According to the English Wikipedia article "Lambda", the Unicode Consortium adopted that spelling because of "preferences expressed by the Greek National Body."

13   Watching the discussion in the python-dev mailing list I thought one reason why `else` was rejected was the lack of consensus on how to indent it within `match`: should `else` be indented at the same level as `match`, or at the same level as `case`?

14   See slide 21 in "Python is Awesome".